

BIG DATA Y COMPUTACIÓN DE ALTO DESEMPEÑO: UNA INTRODUCCIÓN PARA EL ESTADÍSTICO

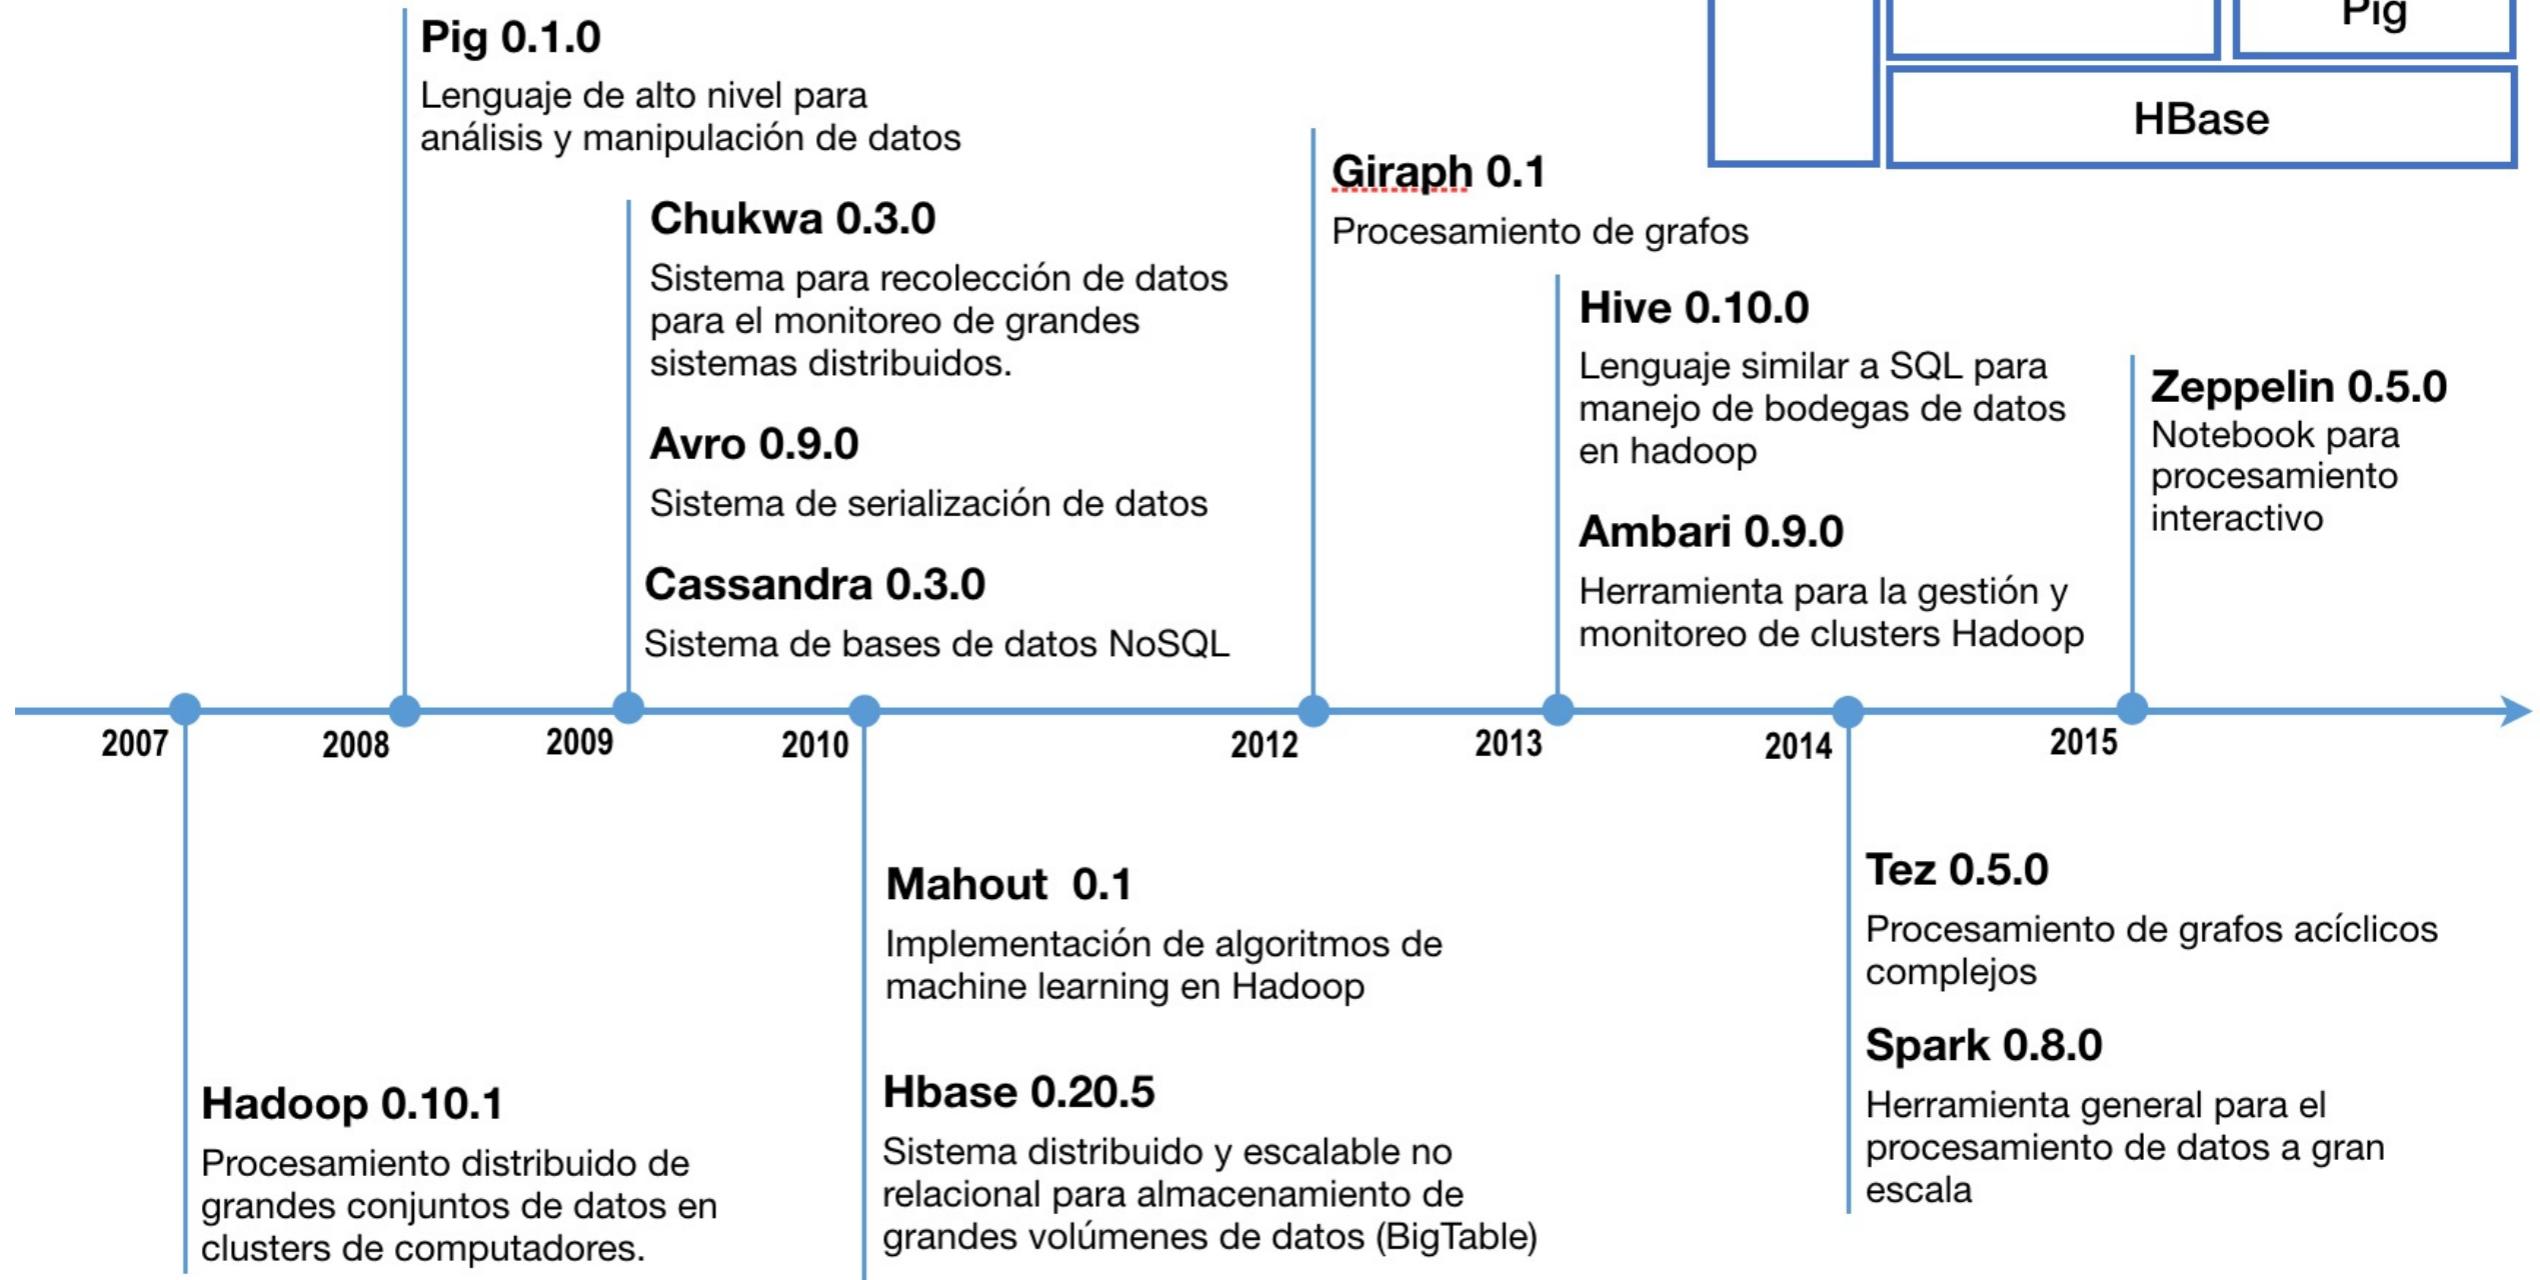
JUAN DAVID VELÁSQUEZ HENAO, MSc, PhD

Profesor Titular

Departamento de Ciencias de la Computación y la Decisión
Facultad de Minas
Universidad Nacional de Colombia, Sede Medellín

-  jdvelasq@unal.edu.co
-  @jdvelasquezh
-  <https://github.com/jdvelasq>
-  <https://goo.gl/prkjAq>
-  <https://goo.gl/vXH8jy>

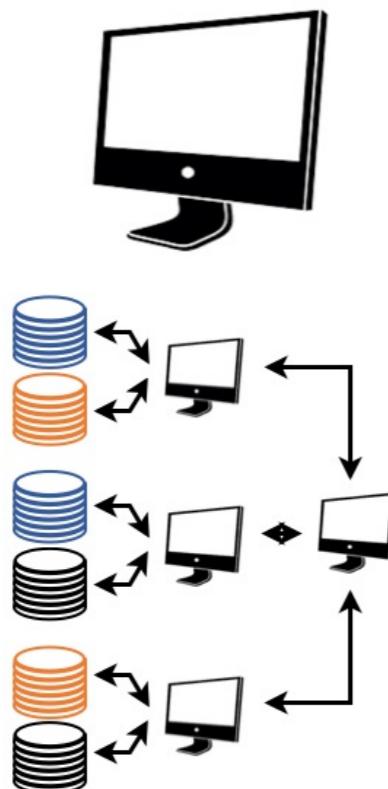
Evolución del Big Data



Servicios Web (2002)

Computación local

Servidores + red + clientes



Cloud computing / utility computing

Servidores y almacenamiento en la nube + internet + clientes locales

Software as a Service (SaaS)

Software almacenado en máquinas suministradas por un tercero.

Aplicaciones accesadas vía un cliente o la Web.

Orientado a aplicaciones de usuario final.

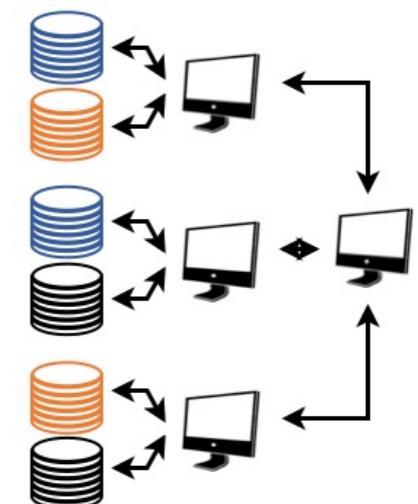
Platform as a Service (PaaS)

Orientado a desarrolladores.

Ambiente de desarrollo gestionado por un tercero.

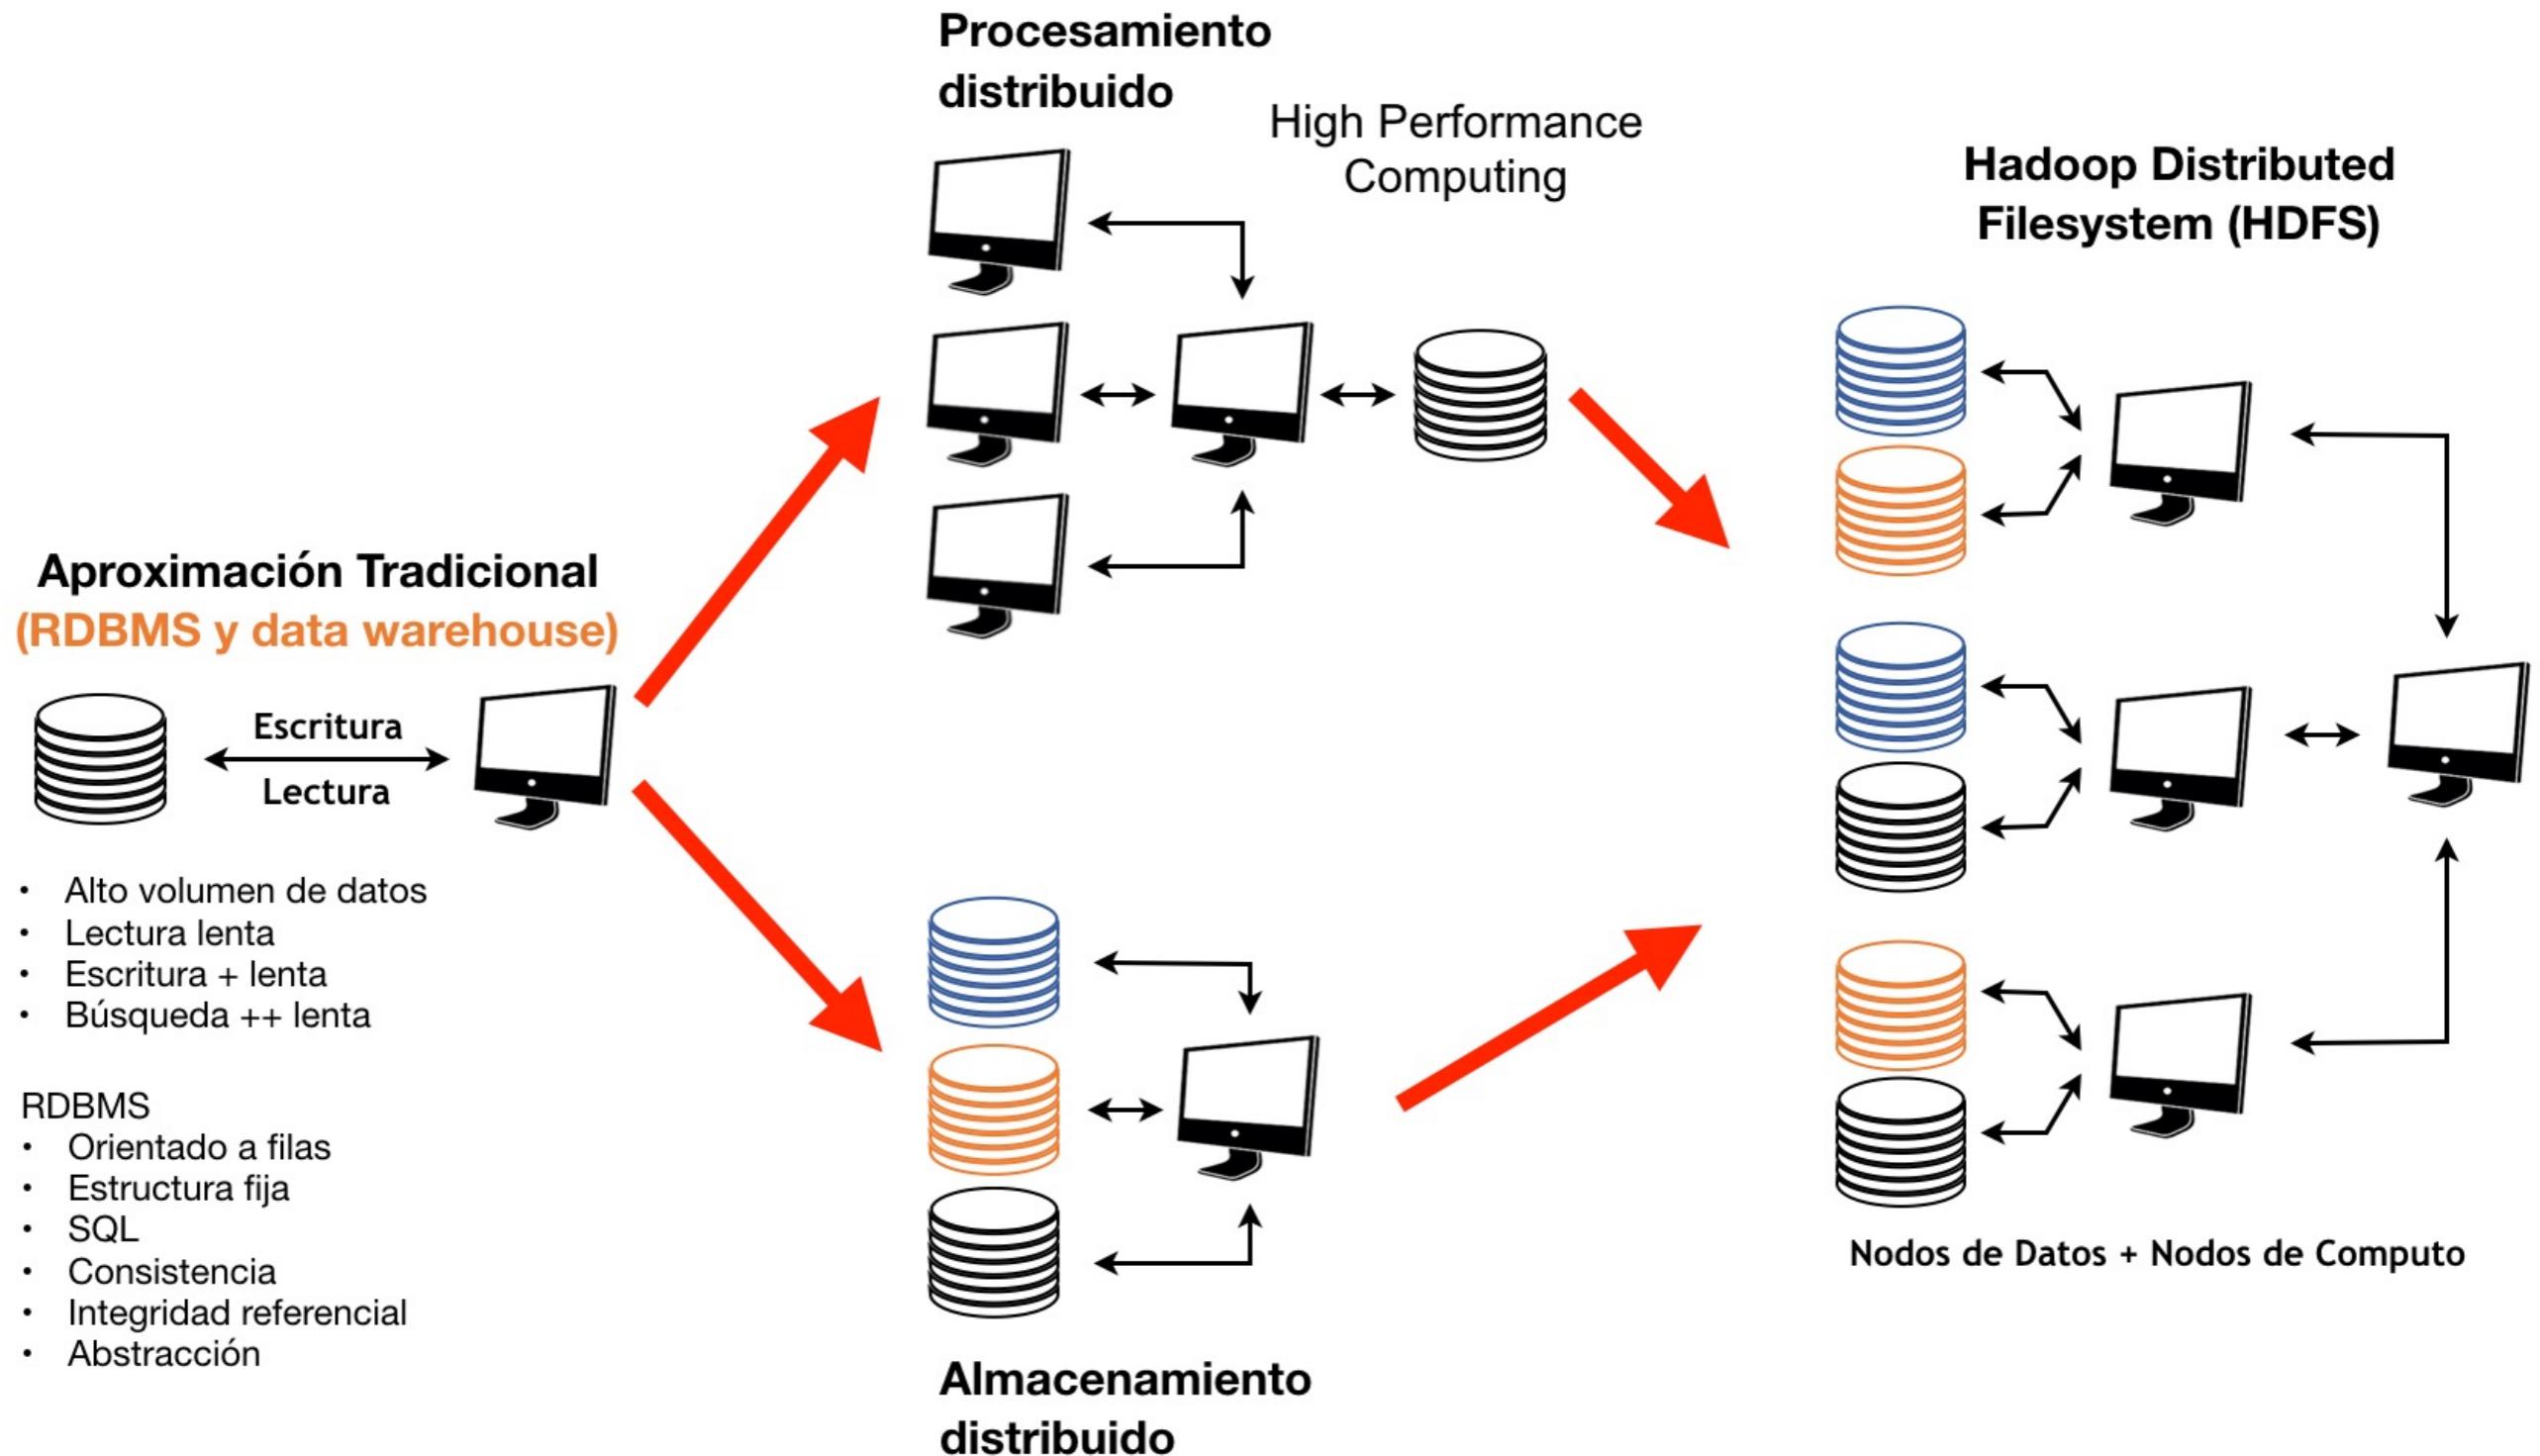
Infrastructure as a Service (IaaS)

Bloques básicos para construcción de ambientes manejados por un tercero
Capacidad de procesamiento, almacenamiento, conectividad, seguridad, etc.

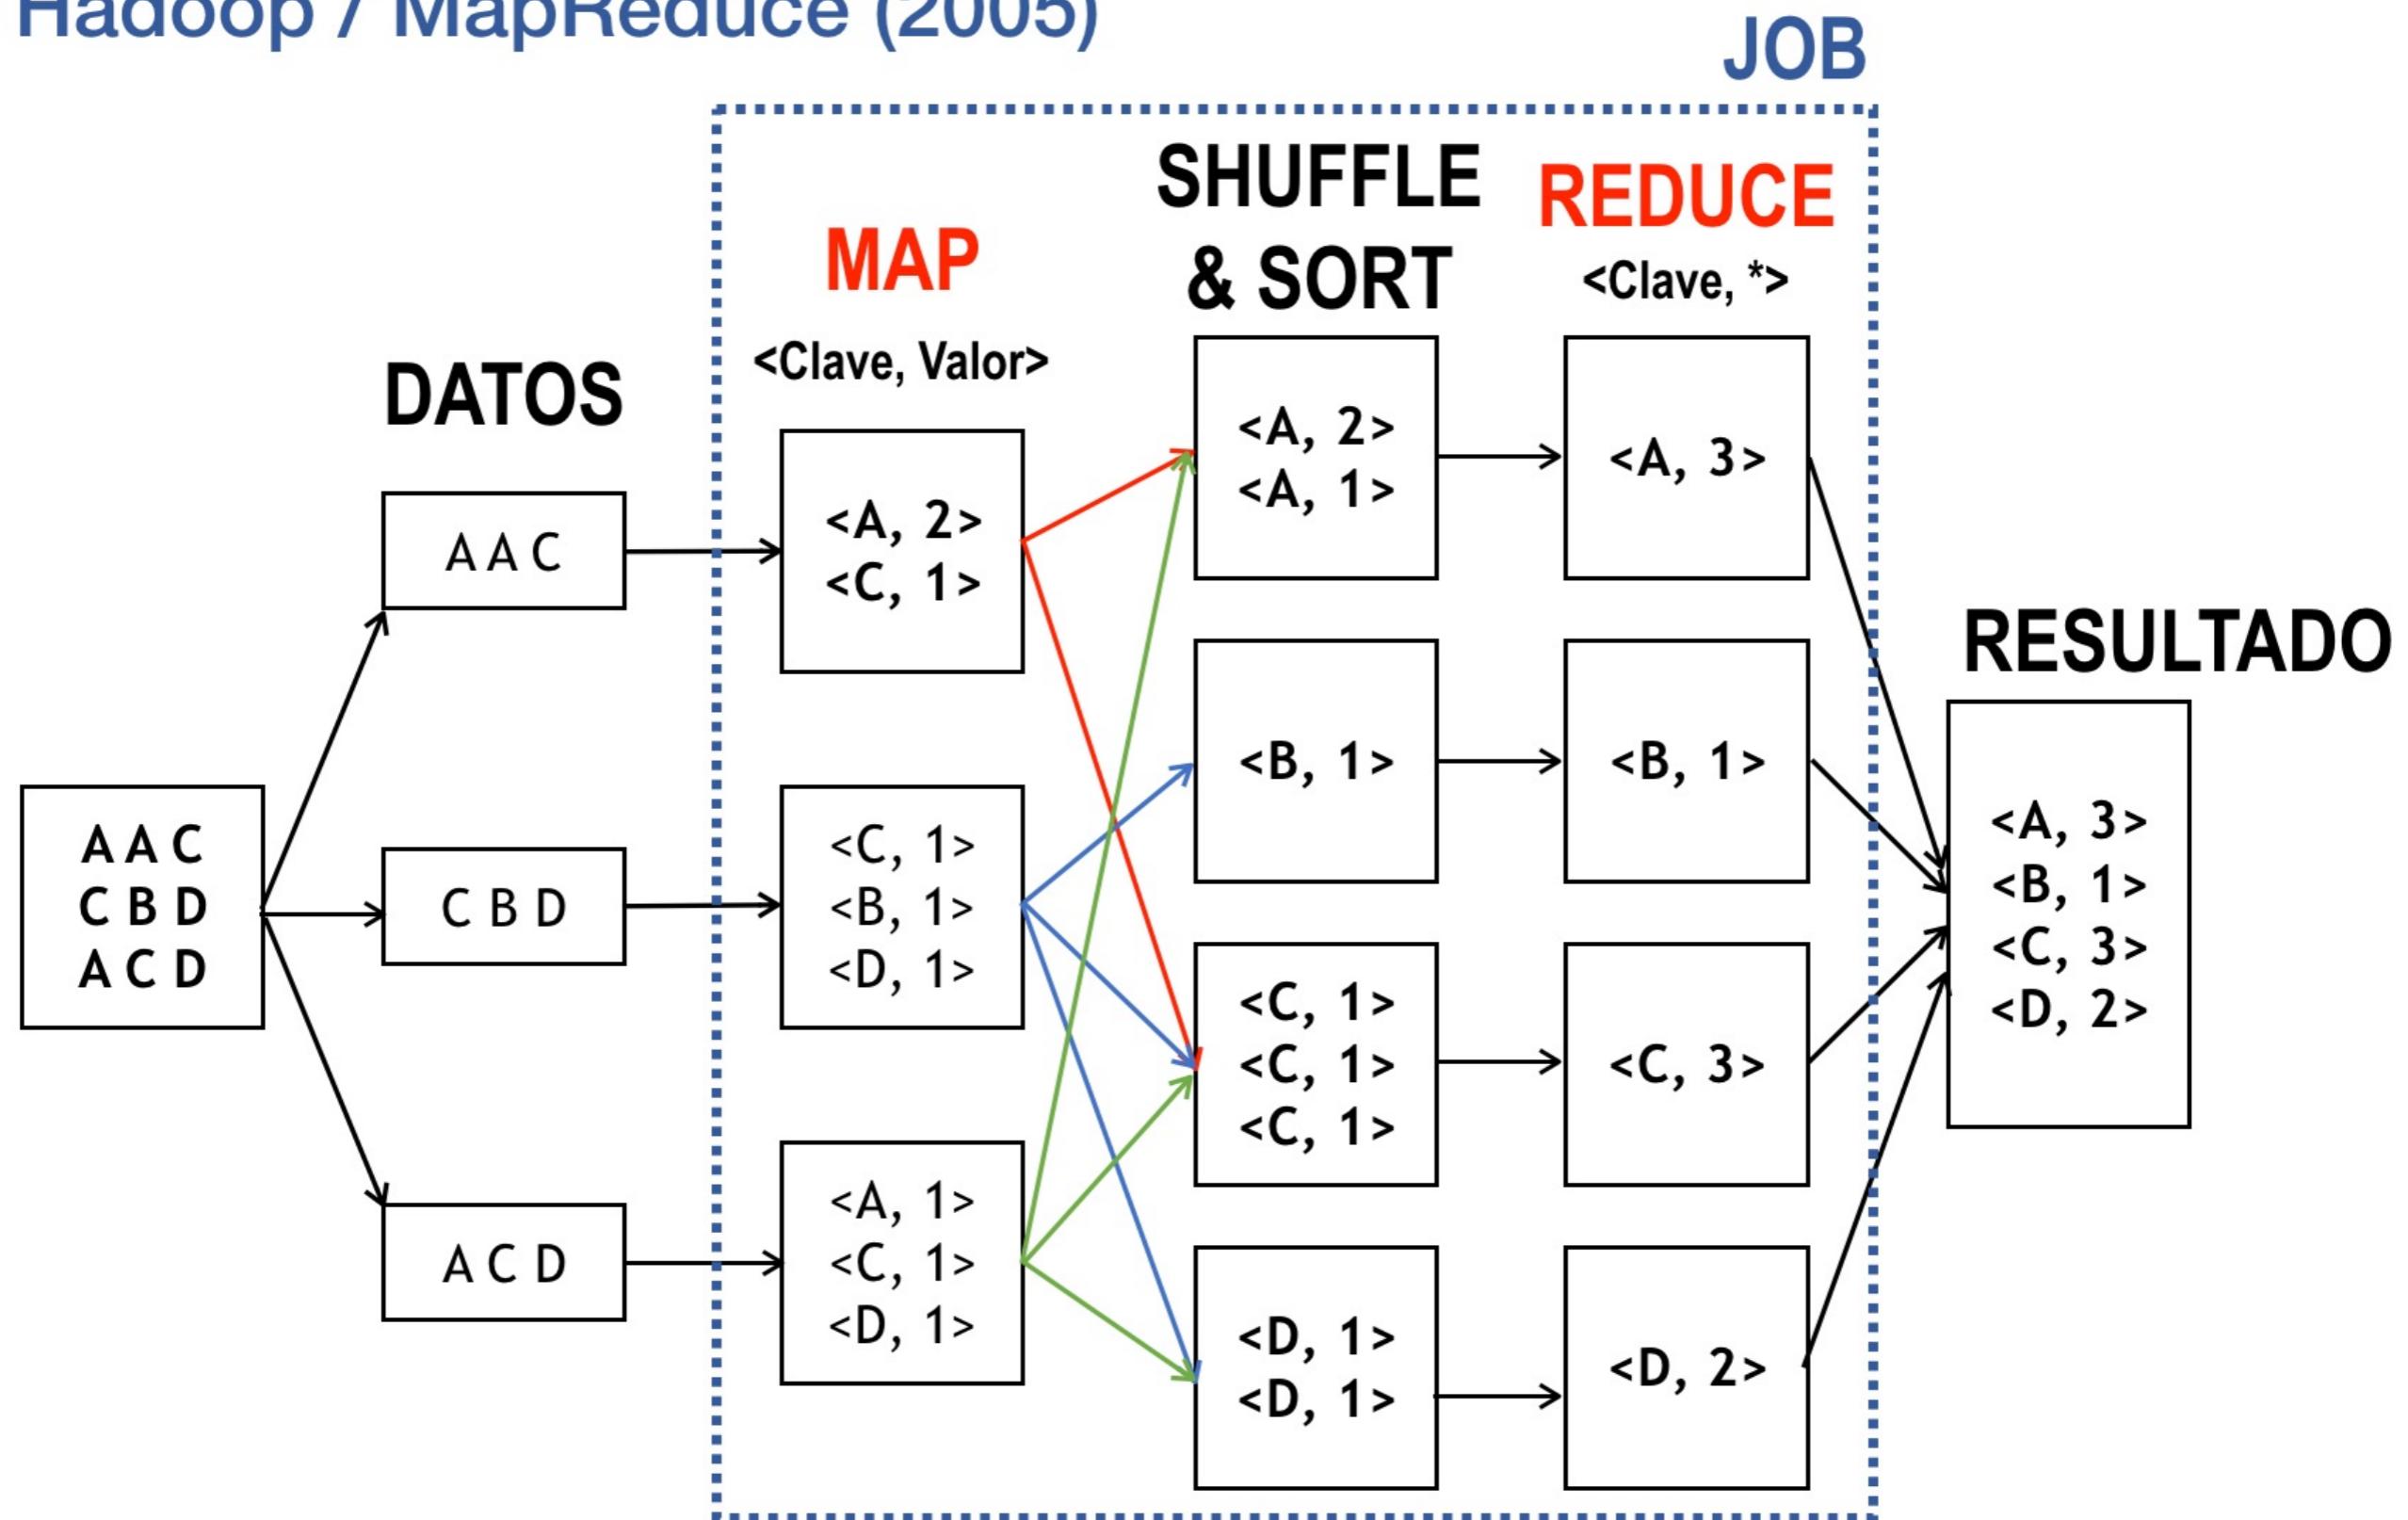


Apache Hadoop / MapReduce

Hadoop / MapReduce



Hadoop / MapReduce (2005)



Generalidades

Modos de ejecución

- Modo local (un solo hilo de ejecución en la máquina local, para aprendizaje y desarrollo). El sistema local de archivos funciona como el HDFS.
- Modo seudo-distribuido (múltiples hilos en la máquina local, para desarrollo). El HDFS es alcanzable como un servicio.
- Cluster (ambiente productivo)

Desarrollo de aplicaciones

- Programación directa en Java (Hadoop)
- Programación usando un lenguaje interpretado (Hadoop-streaming)

Sistemas Operativos

- Microsoft Windows (difícil).
- Mac OS: instalación directa.
- Ubuntu Linux: instalación directa.

Conteo de palabras usando Hadoop-Streaming

Datos

```
In [93]: %%writefile input/text0.txt  
A B C A  
A D D A  
A K M C
```

Overwriting input/text0.txt

```
In [94]: %%writefile input/text1.txt  
B A C Y B  
U O Y Y A  
A B I T
```

Overwriting input/text1.txt

```
In [95]: %%writefile input/text2.txt  
A C D A  
A K B  
A N H I D A
```

Overwriting input/text2.txt

Conteo de palabras usando Hadoop-Streaming

mapper.R

```
In [96]: ##writefile mapper.R  
#!/usr/bin/env Rscript  
  
input <- file('stdin', 'r')  
while(TRUE) {  
    row <- readLines(input, n=1)  
    if( length(row) == 0 ){  
        break  
    }  
    words <- strsplit(row, " ")[[1]]  
    for(word in words){  
        if(word != '')  
            write(cat(word, '\t1', sep=''), "")  
    }  
}
```

Overwriting mapper.R

```
In [119]: ## El programa anterior se hace ejecutable  
!chmod +x mapper.R  
  
## Verificación  
!cat ./input/text*.txt | ./mapper.R | head
```

A	1
B	1
C	1
A	1
A	1
D	1
D	1
A	1
A	1
K	1

Conteo de palabras usando Hadoop-Streaming

reducer.R

```
In [98]: ##writefile reducer.R
#!/usr/bin/env Rscript

curkey <- NULL
total <- 0
input <- file('stdin', 'r')
while(TRUE) {
  row <- readLines(input, n=1)
  if( length(row) == 0 ){
    break
  }
  x <- strsplit(row, "\t")[[1]]
  key <- x[1]
  value <- strtoi(x[2])
  if(!is.null(curkey) && key == curkey){
    total <- total + value
  }
  else{
    if( !is.null(curkey) ) {
      write(cat(curkey, '\t', total), "")
    }
    curkey <- key
    total <- value
  }
}
write(cat(curkey, '\t', total), "")
```

Overwriting reducer.R

```
In [99]: chmod +x reducer.R
!cat ./input/text*.txt | ./mapper.R | sort | ./reducer.R

Warning message:
In readLines(input, n = 1) : incomplete final line found on 'stdin'
A      13
B       5
C       4
D       4
H       1
I       2
K       2
M       1
N       1
O       1
T       1
U       1
Y       3
```

Conteo de palabras usando Hadoop-Streaming

```
In [115]: !hadoop fs -mkdir /user  
!hadoop fs -mkdir /user/jdvelasq  
!hadoop fs -mkdir /user/jdvelasq/input  
!hadoop fs -copyFromLocal input/* /user/jdvelasq/input  
!hadoop fs -ls /user/jdvelasq/input/*  
  
-rw-r--r-- 1 jdvelasq supergroup 25 2018-11-08 23:36 /user/jdvelasq/input/text0.txt  
-rw-r--r-- 1 jdvelasq supergroup 29 2018-11-08 23:36 /user/jdvelasq/input/text1.txt  
-rw-r--r-- 1 jdvelasq supergroup 28 2018-11-08 23:36 /user/jdvelasq/input/text2.txt  
  
In [116]: !hadoop jar $HADOOP_HOME/share/hadoop/tools/lib/hadoop-streaming*.jar \  
-input input -output output -mapper mapper.R -reducer reducer.R  
  
In [117]: !hadoop fs -ls /user/jdvelasq/output  
  
Found 2 items  
-rw-r--r-- 1 jdvelasq supergroup 0 2018-11-08 23:37 /user/jdvelasq/output/_SUCCESS  
-rw-r--r-- 1 jdvelasq supergroup 79 2018-11-08 23:37 /user/jdvelasq/output/part-00000  
  
In [118]: !hadoop fs -cat /user/jdvelasq/output/part-00000  
  
A 13  
B 5  
C 4
```

Conteo de palabras usando Hadoop-Streaming

Hadoop Overview Datanodes Datanode Volume Failures Snapshot Startup Progress Utilities ▾

Browse Directory

/

Show 25 entries Search:

Permission	Owner	Group	Size	Last Modified	Replication	Block Size	Name
drwxr-xr-x	jdvelasq	supergroup	0 B	Nov 09 07:37	0	0 B	Volumes
drwxr-xr-x	jdvelasq	supergroup	0 B	Nov 08 23:31	0	0 B	input
drwxr-xr-x	jdvelasq	supergroup	0 B	Nov 09 07:36	0	0 B	tmp
drwxr-xr-x	jdvelasq						

Showing 1 to 4 of 4 entries

Hadoop, 2018.

A red arrow points from the "Logs" link in the Utilities dropdown menu to the "Logs" link in the "Browse the file system" box.

Hadoop Overview Datanodes Datanode Volume Failures Snapshot Startup Progress Utilities ▾

Browse Directory

/user/jdvelasq/output

Show 25 entries Search:

Permission	Owner	Group	Size	Last Modified	Replication	Block Size	Name
-rw-r--r--	jdvelasq	supergroup	0 B	Nov 08 23:37	1	128 MB	_SUCCESS
-rw-r--r--	jdvelasq	supergroup	79 B	Nov 08 23:37	1	128 MB	part-00000

Showing 1 to 2 of 2 entries

Previous 1 Next

Hadoop, 2018.

Problemas típicos en los que puede usarse Map/Reduce

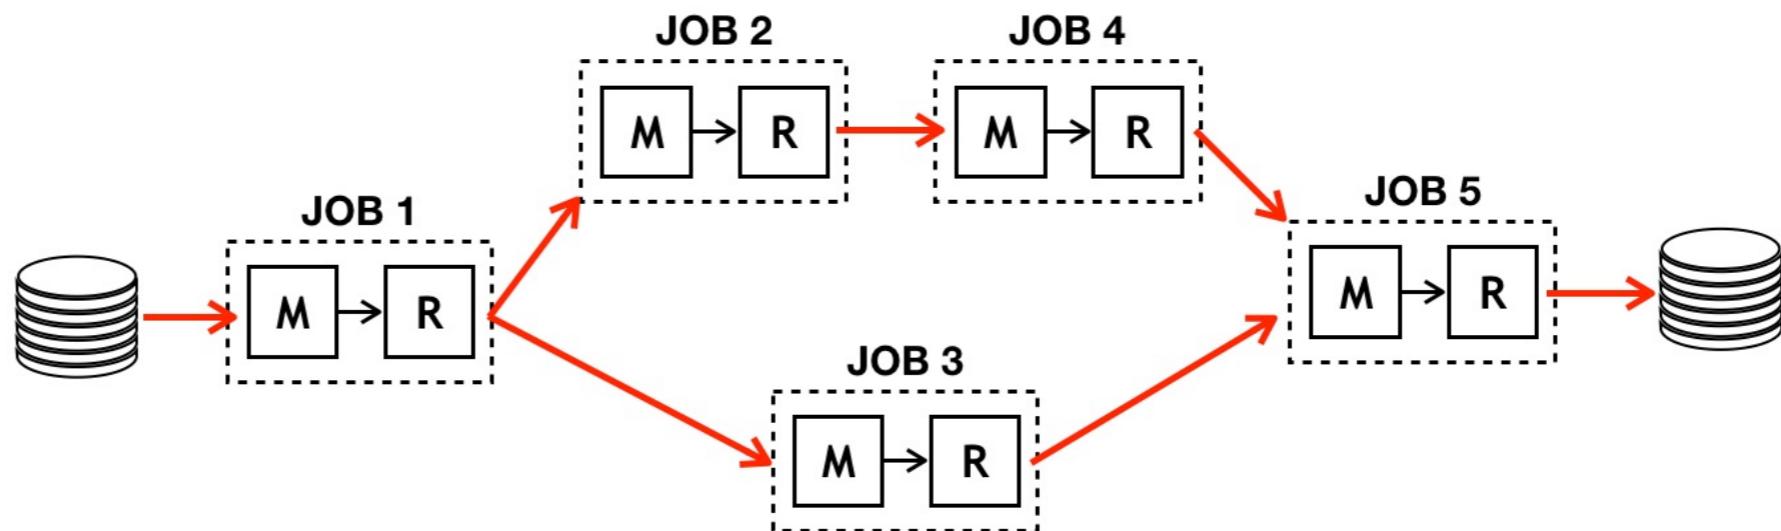
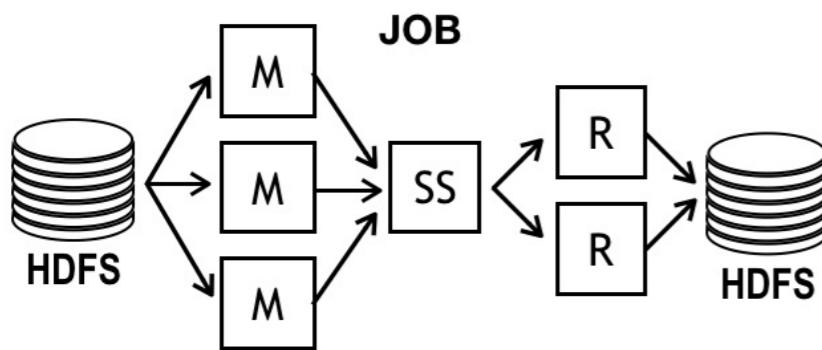
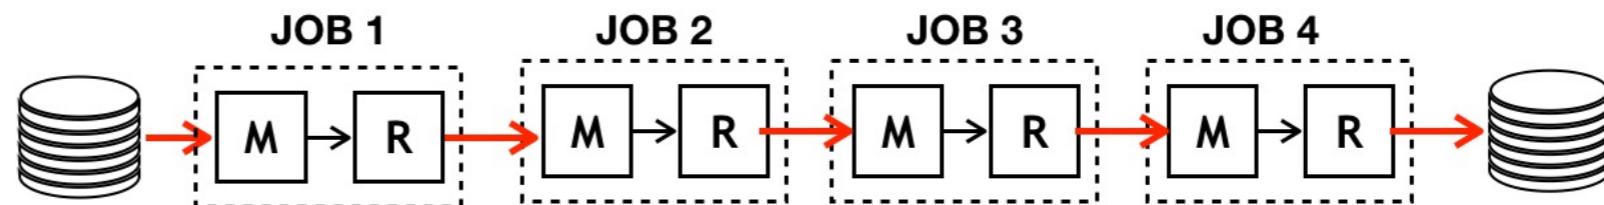
- Valores máximo o mínimo para una clave.
- Ordenamiento del archivo.
- Sumas y promedios por claves.
- Obtener los N registros más ...
- Asociar por claves:

0	A, B	=====>	A 0, 1
1	A, B		B 0, 1, 2
2	B, C		C 2

La mayor parte de problemas no pueden solucionarse usando un proceso simple Map/Reduce



Jobs

Hadoop / MapReduce (2005)



La coordinación entre jobs debe realizarse explícitamente en la consola de comandos

Apache Tez



Apache Pig

Conteo de palabras usando Pig

```
In [1]: %load_ext bigdata  
%pig_init
```

```
In [5]: %%pig  
lines = LOAD 'input/text*.txt' AS (line:CHARARRAY);  
  
-- genera una tabla llamada words con una palabra por registro  
words = FOREACH lines GENERATE FLATTEN(TOKENIZE(line)) AS word;  
  
-- agrupa los registros que tienen la misma palabra  
grouped = GROUP words BY word;  
  
-- genera una variable que cuenta las ocurrencias por cada grupo  
wordcount = FOREACH grouped GENERATE group, COUNT(words);  
  
-- selecciona las primeras 15 palabras  
s = LIMIT wordcount 15;  
  
-- imprime en pantalla las primeras 15 palabras  
STORE s INTO 'output-pig';
```

```
In [6]: !hadoop fs -ls output-pig/
```

```
Found 2 items  
-rw-r--r--    1 jdvelasq supergroup          0 2018-11-09 07:22 output-pig/_SUCCESS  
-rw-r--r--    1 jdvelasq supergroup      53 2018-11-09 07:22 output-pig/part-r-00000
```

```
In [7]: !hadoop fs -cat /user/jdvelasq/output-pig/part-r-00000
```

A	13
B	5
C	4

Conteo de palabras usando Pig

Browse Directory

/user/jdvelasq

Go!



Show 25 ↑ entries

Search:

File information - part-r-00000

Download Head the file (first 32K) Tail the file (last 32K)

Showing 1 to 3 of 3

Hadoop, 2018

Block information -- Block 0

Block ID: 1073741847
Block Pool ID: BP-1224940242-192.168.1.58-1541734190285
Generation Stamp: 1023
Size: 53
Availability:

- 192.168.1.58

File contents

A 13
B 5
C 4
D 4
H 1
I 2
K 2
M 1

Replication	Block Size	Name	
0	0 B	input	trash
0	0 B	output	trash
0	0 B	output-pig	trash

Previous 1 Next

<http://nbviewer.jupyter.org/github/jdvelasq/AGD-03-Pig/blob/master/02-basics.ipynb>

<http://nbviewer.jupyter.org/github/jdvelasq/AGD-03-Pig/blob/master/04-tipos-de-datos.ipynb>

Apache Hive

Conteo de palabras

```
In [8]: %hive_init
```

```
Hive initialized!
```

```
In [9]: %%hive
```

```
DROP TABLE IF EXISTS docs;
DROP TABLE IF EXISTS word_counts;

CREATE TABLE docs (line STRING);

LOAD DATA LOCAL INPATH
    'input/text*.txt'
OVERWRITE INTO TABLE docs;

CREATE TABLE word_counts
AS
    SELECT word, count(1) AS count
    FROM
        (SELECT explode(split(line, '\\s')) AS word FROM docs) w
GROUP BY
    word
ORDER BY
    word;

SELECT * FROM word_counts LIMIT 10;
```

	7
A	13
B	5
C	4
D	4
H	1
I	2
K	2
M	1
N	1

Características de Apache Hive

- El lenguaje tiene muchas similitudes con SQL y resulta fácil
- Los archivos se guardan directamente en disco duro como texto.
- A continuación se presentan los principales puntos de diferencia con otros sistemas de bases de datos.
- Obtener los N registros más

Hive

In [12]:

```
%%hive
DROP TABLE IF EXISTS tbl0;
CREATE TABLE tbl0 (
    c1 INT,
    c2 STRING,
    c3 INT,
    c4 DATE,
    c5 ARRAY<CHAR(1)>,
    c6 MAP<STRING, INT>
)
ROW FORMAT DELIMITED
FIELDS TERMINATED BY ','
COLLECTION ITEMS TERMINATED BY ':'
MAP KEYS TERMINATED BY '#'
LINES TERMINATED BY '\n';

LOAD DATA LOCAL INPATH 'tbl0.csv' INTO TABLE tbl0;

SELECT * FROM tbl0;
```

c1	c2	c3	c4	c5	c6
1	D	4	2016-06-25	["a", "e", "c", "d"]	{"bb":10, "dd":20, "cc":40}
2	C	4	2015-05-14	["a", "c"]	{"dd":40, "bb":20, "cc":10}
3	D	6	2014-12-26	["b", "d"]	{"aa":10, "bb":40}
4	D	5	2016-06-25	["a", "c", "e", "b", "d"]	{"bb":40, "dd":20, "aa":10, "cc":30}
5	C	7	2016-05-23	["d", "e", "c"]	{"cc":10, "aa":20}
6	A	2	2018-06-14	["a", "d"]	{"aa":20}
7	B	3	2014-12-22	["a", "e", "d"]	{"cc":20}
8	C	6	2015-08-20	["d", "a", "c", "e"]	{"cc":10, "aa":20}
9	B	3	2017-12-08	["b", "a", "c", "e"]	{"cc":40, "dd":10, "aa":30, "bb":20}
10	B	7	2015-07-28	["c", "d", "e", "a", "b"]	{"aa":10, "dd":40, "cc":30}

In [11]:

```
%%writefile tbl0.csv
1,D,4,2016-06-25,a:e:c:d,bb#10:dd#20:cc#40
2,C,4,2015-05-14,a:c,dd#40:bb#20:cc#10
3,D,6,2014-12-26,b:d,aa#10:bb#40
4,D,5,2016-06-25,a:c:e:b:d,bb#40:dd#20:aa#10:cc#30
5,C,7,2016-05-23,d:e:c,cc#10:aa#20
6,A,2,2018-06-14,a:d,aa#20
7,B,3,2014-12-22,a:e:d,cc#20
8,C,6,2015-08-20,d:a:c:e,cc#10:aa#20
9,B,3,2017-12-08,b:a:c:e,cc#40:dd#10:aa#30:bb#20
10,B,7,2015-07-28,c:d:e:a:b,aa#10:dd#40:cc#30
```

Writing tbl0.csv

Apache Spark

Conteo de palabras

```
In [1]: import findspark
findspark.init()

from pyspark import SparkConf, SparkContext
from operator import add

APP_NAME = "My First Spark Application"

def tokenize(text):
    return text.split()

def main(sc):
    text = sc.textFile('input/text*.txt')
    words = text.flatMap(tokenize)
    wc = words.map(lambda x: (x,1))
    counts = wc.reduceByKey(add)
    counts.saveAsTextFile("output-spark")

if __name__ == "__main__":
    conf = SparkConf().setAppName(APP_NAME)
    conf = conf.setMaster("local[*]")
    sc = SparkContext(conf=conf)
    main(sc)
```

```
In [2]: !hadoop fs -ls output-spark
```

```
Found 4 items
-rw-r--r-- 1 jdvelasq supergroup      0 2018-11-09 08:39 output-spark/_SUCCESS
-rw-r--r-- 1 jdvelasq supergroup    45 2018-11-09 08:39 output-spark/part-00000
-rw-r--r-- 1 jdvelasq supergroup    37 2018-11-09 08:39 output-spark/part-00001
-rw-r--r-- 1 jdvelasq supergroup    36 2018-11-09 08:39 output-spark/part-00002
```

```
In [6]: !hadoop fs -cat /user/jdvelasq/output-spark/part-00000
('B', 5)
('C', 4)
('D', 4)
('T', 1)
('N', 1)
```

Aplicación arbitraria de funciones

```
In [7]: ## en este ejemplo se pasa una función arbitraria a `map`  
from operator import add  
rdd = sc.textFile('input/text*.txt')  
rdd = rdd.map(len)  
print(rdd.collect())  
rdd = rdd.reduce(add)  
rdd
```

```
[7, 7, 7, 1, 9, 9, 7, 1, 7, 6, 11, 1]
```

```
Out[7]: 73
```

Spark QL

```
In [7]: from pyspark.sql import Row
```

```
## crea el contexto  
sc = spark.sparkContext
```

```
## lee el archivo como lineas de texto  
rdd = sc.textFile("people.csv")
```

```
## parte por las comas  
rdd = rdd.map(lambda row: row.split(","))
```

```
## cambia los tipos de los datos  
rdd = rdd.map(lambda p: Row(id=p[0],  
                           firstname=p[1],  
                           surname=p[2],  
                           birthdate=p[3],  
                           color=p[4],  
                           quantity=int(p[5])))
```

```
## crea el DataFrame  
df = spark.createDataFrame(rdd)  
df.createOrReplaceTempView("df")  
df.show()
```

	birthdate	color	firstname	id	quantity	surname
	1971-07-08	green	Vivian	1		
	1974-05-23	green	Karen	2		
	1973-04-22	orange	Cody	3		
	1975-01-29	black	Roth	4		
	1974-07-03	blue	Zoe	5		

```
In [ ]: query = """  
SELECT DISTINCT  
    _C5  
FROM  
    csv.`datos.csv`  
ORDER BY  
    _c5  
"""  
  
spark.sql(query).write.save('temp', format="csv")
```

```
In [13]: ## Filtrado
```

```
df.filter(df['color'] == 'blue').show()
```

	id	firstname	surname	birthdate	color	quantity
	5	Zoe	Conway	1974-07-03 00:00:00	blue	2
	7	Driscoll	Klein	1970-10-05 00:00:00	blue	5
	15	Hope	Silva	1970-07-01 00:00:00	blue	5

MLlib: Main Guide

- Basic statistics
- Data sources
- Pipelines
- Extracting, transforming and selecting features
- Classification and Regression
- Clustering
- Collaborative filtering
- Frequent Pattern Mining
- Model selection and tuning
- Advanced topics

MLlib: RDD-based API Guide

- Data types
- Basic statistics
- Classification and regression
- Collaborative filtering
- Clustering
- Dimensionality reduction
- Feature extraction and transformation
- Frequent pattern mining
- Evaluation metrics
- PMML model export
- Optimization (developer)

Machine Learning Library (MLlib) Guide

MLlib is Spark's machine learning (ML) library. Its goal is to make practical machine learning scalable and efficient. It provides tools such as:

- ML Algorithms: common learning algorithms such as classification, regression, clustering, and dimensionality reduction.
- Featurization: feature extraction, transformation, dimensionality reduction, and selection.
- Pipelines: tools for constructing, evaluating, and tuning ML Pipelines.
- Persistence: saving and load algorithms, models, and Pipelines.
- Utilities: linear algebra, statistics, data handling, etc.

Announcement: DataFrame-based API is prioritized

The MLlib RDD-based API is now in maintenance mode.

As of Spark 2.0, the [RDD-based APIs](#) in the `spark.mllib` package have entered maintenance mode. The Machine Learning API for Spark is now the [DataFrame-based API](#) in the `spark.ml` package.

What are the implications?

- MLlib will still support the RDD-based API in `spark.mllib` with bug fixes.
- MLlib will not add new features to the RDD-based API.
- In the Spark 2.x releases, MLlib will add features to the DataFrames-based API to reach feature parity.
- After reaching feature parity (roughly estimated for Spark 2.3), the RDD-based API will be deprecated.
- The RDD-based API is expected to be removed in Spark 3.0.

Why is MLlib switching to the DataFrame-based API?

- DataFrames provide a more user-friendly API than RDDs. The many benefits of DataFrames include SQL/DataFrame queries, Tungsten and Catalyst optimizations, and uniform APIs across languages.
- The DataFrame-based API for MLlib provides a uniform API across ML algorithms and across languages.
- DataFrames facilitate practical ML Pipelines, particularly feature transformations. See the [Pipelines guide](#).

What is "Spark ML"?

* "Spark ML" is not an official name but occasionally used to refer to the MLlib DataFrame-based API.

Scala Java Python R

In SparkR we provide Spark SQL data source API for loading image data as DataFrame.

```
> df = read.df("data/mllib/images/origin/kittens", "image")
> head(select(df, df$image.origin, df$image.width, df$image.height))

1      file:///spark/data/mllib/images/origin/kittens/54893.jpg
2      file:///spark/data/mllib/images/origin/kittens/DP802813.jpg
3 file:///spark/data/mllib/images/origin/kittens/29.5.a_b_EGDP022204.jpg
4      file:///spark/data/mllib/images/origin/kittens/DP153539.jpg
width height
1  300    311
2  199    313
3  300    200
4  300    296
```

Scala Java Python R

More details on parameters can be found in the [R API documentation](#).

```
# Load training data
df <- read.df("data/mllib/sample_libsvm_data.txt", source = "libsvm")
training <- df
test <- df

# Fit an binomial logistic regression model with spark.logit
model <- spark.logit(training, label ~ features, maxIter = 10, regParam = 0.3, elasticNetParam = 0.8)

# Model summary
summary(model)

# Prediction
predictions <- predict(model, test)
head(predictions)
```

[Using sparklyr](#)[Configuring connections](#)[Troubleshooting](#)[Guides](#)[Manipulating data](#)[Machine Learning](#)[Understanding Caching](#)[Deployment Options](#)[Distributed R](#)[Data Lakes](#)[ML Pipelines](#)[Text mining](#)[Stream Analysis](#)[Extend sparklyr](#)[Using H2O](#)[Graph Analysis](#)[Production pipelines](#)

sparklyr: R interface for Apache Spark

[build](#) passing[CRAN](#) 0.9.2[codecov](#) 80%[chat](#) on gitter

- Connect to [Spark](#) from R. The sparklyr package provides a complete [dplyr](#) backend.
- Filter and aggregate Spark datasets then bring them into R for analysis and visualization.
- Use Spark's distributed [machine learning](#) library from R.
- Create [extensions](#) that call the full Spark API and provide interfaces to Spark packages.



Installation

You can install the **sparklyr** package from CRAN as follows:

```
install.packages("sparklyr")
```

You should also install a local version of Spark for development purposes:

```
library(sparklyr)  
spark_install(version = "2.1.0")
```

Apache Zeppelin



demo-python




```
print('hola mund0')
```

FINISHED

```
hola mund0
```

READY



Interpreters

[Overview](#)[Spark](#)[JDBC](#)[Python](#)[Alluxio](#)[Beam](#)[BigQuery](#)[Cassandra](#)[Elasticsearch](#)[Flink](#)[Geode](#)[Groovy](#)[HBase](#)[HDFS](#)[Hive](#)[Ignite](#)[Kylin](#)

Apache DRILL

```
SELECT * FROM dfs.`<path-to-installation>/apache-drill-<version>/sample-data/region.parquet`;
```

R_REGIONKEY	R_NAME	R_COMMENT
0	AFRICA	lar deposits. blithe
1	AMERICA	hs use ironic, even
2	ASIA	ges. thinly even pin
3	EUROPE	ly final courts cajo
4	MIDDLE EAST	uickly special accou

5 rows selected (0.409 seconds)

BIG DATA Y COMPUTACIÓN DE ALTO DESEMPEÑO: UNA INTRODUCCIÓN PARA EL ESTADÍSTICO

JUAN DAVID VELÁSQUEZ HENAO, MSc, PhD

Profesor Titular

Departamento de Ciencias de la Computación y la Decisión
Facultad de Minas
Universidad Nacional de Colombia, Sede Medellín

-  jdvelasq@unal.edu.co
-  @jdvelasquezh
-  <https://github.com/jdvelasq>
-  <https://goo.gl/prkjAq>
-  <https://goo.gl/vXH8jy>